Tasks:

Create an intelligent agent to learn how to play Pong, using Q-Learning.
The agent should play against:

1. A random strategy: at each frame, chooses a random direction (up/down) to move.
2. A greedy strategy: always chooses the best action.
3. An epsilon-greedy strategy: while unexplored actions exist, choose one randomly. If all actions are explored, with a probability of epsilon, choose a random action and with 1-epsilon, choose the best action.

Readme:

Run the source:

python Pong.py

At start, the agent learns how to play. (it may require many, many training games to learn).
Then, using the model learned, the agent plays against a random adversary, a greedy one and an epsilon-greedy one.